Architectural Foundations for High-Throughput Holographic Neural Weights via MHz-Spooled Image Streams

Introduction: The Paradigm Shift in Neural Weight Instantiation

The prevailing architecture of artificial neural networks rests upon the von Neumann computational model, wherein the synaptic weight is defined as a discrete, mathematically based functional relationship—a numerical shard—connecting localized nodes within a computational graph. These weights are instantiated as floating-point values in memory, requiring massive data transport to Arithmetic Logic Units (ALUs) for tensor operations, typically matrix-vector multiplications. As the dimensional complexity of models scales, the latency and energy costs associated with fetching these localized shards degrade system throughput precipitously.

To circumvent the fundamental limitations of localized tensor mathematics, a radical architectural departure is required: the abandonment of the mathematical shard in favor of a byte-by-byte processable physical hologram. In this paradigm, a neural weight is not a discrete numerical value, but rather a dynamic, physical interference pattern. This approach shifts the computational burden entirely from floating-point arithmetic to data streaming bandwidth. By leveraging industry-standard image formats—such as the Tagged Image File Format (TIFF) for high-dynamic-range phase maps and the Graphics Interchange Format (GIF) for rapid, lossless Lempel-Ziv-Welch (LZW) streaming—pipelined via Field Programmable Gate Arrays (FPGAs) at Megahertz (MHz) refresh rates, it becomes possible to instantiate a continuously streaming optical intelligence engine.

This report presents an exhaustive theoretical, optical, and architectural analysis of constructing high-throughput holographic neural weights. The analysis synthesizes historical precedents in holographic neural technology, Vector Symbolic Architectures (VSA), Fourier optics, 5D physical storage mediums, high-frequency Acousto-Optic Deflectors (AODs), and direct memory access byte-streaming pipelines.

The Theoretical Substrate: Beyond Tensor Shards

The ambition to transcend discrete mathematical shards necessitates a representational space capable of continuous, distributed analogical encoding. In localized networks, the loss or corruption of a single critical weight drastically alters the output. In a holographic system, information is distributed continuously across the entire representational substrate, ensuring robustness and enabling quantum-logical phase interactions.

Holographic Neural Technology and Complex Phase Space

The foundational premise of non-mathematical weight relationships originates in Holographic Neural Technology (HNeT), pioneered by John Sutherland[1]. Traditional neural models represent magnitude and certainty as distinct scalar values. HNeT, conversely, encodes features and certainty simultaneously within the phase terms of complex numbers[1]. By mapping the phase orientation of complex numbers to cognitive features, the system operates on principles of phase coherence and decoherence rather than linear summation[2].

This complex phase space mimics the holoflux—a wholistic, indivisible process defined by quantum logic, where coherent, organized behavior emerges across macroscopic observables[3]. In a physical holographic neural network (HNN), this mathematical formalism dictates that the network's global state can be accurately approximated by sampling a small cross-section of the underlying microstate[5]. Theoretical models of computational complexity analyze these as "fuzzy Boolean functions," where a function $f : \{0,1\}^n \to [0,1]$ is deemed "holographic" if its value can be recovered up to a small error, with high probability, from a bounded number of randomly chosen coordinates[5]. This structural property proves that bounded random recoverability is qualitatively equivalent to bounded neural network properties, meaning the physical hologram natively executes complex polynomial approximations simply by existing in space[5].

Vector Symbolic Architectures and Holographic Reduced Representation

To encode structured, hierarchical concepts within this holographic medium, the architecture relies upon Vector Symbolic Architectures (VSA) and Holographic Reduced Representations (HRR)[8]. Developed by Tony Plate, HRR maps symbols to high-dimensional, distributed hypervectors, bypassing the need for discrete associative arrays[8].

HRR operations map probability and logic directly into vector space through three primary mechanisms. First, similarity is determined by the dot product of two vectors, estimating analogical closeness[12]. Second, bundling (superposition) combines sets of symbols into a single memory trace via vector addition, distributing similarity over the bundle[10]. Third, binding associates two vectors to create a novel concept vector of the exact same dimensionality[8].

In HRR, binding is executed via circular convolution, denoted as $\mathbf{z} = \mathbf{x} \circledast \mathbf{y}$ [8]. In traditional electronic hardware, circular convolution is computationally prohibitive for millions of parameters, requiring continuous Fast Fourier Transforms (FFTs)[15]. However, in the frequency domain, circular convolution simplifies to element-wise complex multiplication—multiplying the moduli and adding the phases[10].

HRR Operation	Mathematical Formalism	Optical Equivalent	Cognitive Function
Similarity	Dot Product: $\mathbf{x} \cdot \mathbf{y}$	Interference Intensity	Analogical matching

Bundling	Addition: $\mathbf{x} + \mathbf{y}$	Wavefront Superposition	Memory trace generation
Binding	Circular Convolution: $\mathbf{x} \circledast \mathbf{y}$	$4f$ Lens Fourier Transform	Concept association

Furthermore, continuous data types are encoded using fractional binding to create Spatial Semantic Pointers (SSPs)[12]. By taking the Fourier transform of a base vector and applying a real-valued exponent, continuous spatial trajectories are represented as unitary vectors (vectors whose Fourier components possess a magnitude of one)[12]. This continuous, phase-based representation perfectly aligns with physical optics, where the byte stream defines the phase shift applied to a coherent light source.

Optical Physics: Translating Bytes into Diffractive Wavefronts

Transitioning from the abstract mathematics of HRR to physical space requires the deployment of Diffractive Deep Neural Networks (D2NNs) and Fourier optics. The objective is to replace the Arithmetic Logic Unit with the physical propagation of light[18].

The Mach-Zehnder Limitation vs. Diffractive Propagation

Early optical neural networks attempted to replicate traditional matrix multiplication using Mach-Zehnder Interferometers (MZIs)[20]. An MZI mesh relies on physical waveguides and localized heaters to alter the refractive index of one arm of the interferometer, computing the product between a 2×2 unitary matrix and a 2-element data vector[21]. While faster than electronics, MZIs remain philosophically tethered to the "shard" model: they are rigid, physically discrete, and highly localized components that scale poorly for massive parallelism[24]. Conversely, D2NNs utilize free-space diffractive optical elements (DOEs)[26]. Each pixel on a diffractive layer acts as an artificial neuron, modifying the complex amplitude of the incoming coherent light[28]. The propagation of light between the l-th and $(l+1)$-th layers is modeled by the Rayleigh-Sommerfeld diffraction equation[28]:

$$w_i^l(x, y, z) = \frac{z - z_i}{r^2} \left(\frac{1}{2\pi r} + \frac{1}{j\lambda} \right) e^{j \frac{2\pi r}{\lambda}}$$

where r is the Euclidean distance and λ is the wavelength[28]. The transmission coefficient at each diffractive neuron—the actual "holographic weight"—is defined as:

$$t_i^l(x_i, y_i, z_i) = a_i^l(x_i, y_i, z_i)e^{j\phi_i^l(x_i, y_i, z_i)}$$

where a dictates amplitude modulation and ϕ dictates the phase delay[28]. By cascading these layers, the network learns complex decision boundaries natively in the optical domain[31].

The $4f$ Optical Correlator System

To execute the circular convolution required for HRR binding instantaneously, the architecture employs a $4f$ optical system[27]. A $4f$ system comprises two objective lenses separated by their focal lengths, with a spatial light modulator situated at the intermediate Fourier plane[35]. When light passes through the first lens, it naturally undergoes a continuous two-dimensional spatial Fourier transform[33]. If a byte-streamed phase mask (the hologram) is applied at the Fourier plane, the light is multiplied by the mask's transfer function. The second lens then performs an inverse Fourier transform, projecting the result onto a detector[37]. This physical process perfectly mimics the $O(1)$ circular convolution necessary for HRR vector binding, entirely eliminating algorithmic latency[15].

Recent advancements in D2NN design have introduced ReLU phase-limiting functions to the physical masks. By restricting the phase parameters during modulation, the network suppresses zero-order noise and background diffractive scattering, raising inference accuracies to parity with deep electronic networks while maintaining speed-of-light execution[34].

The Volumetric Storage Medium: 5D Optical Quartz

Before the holographic weights can be streamed at MHz rates to the active modulator, the immense volume of phase data must be stored physically. Traditional magnetic or solid-state drives struggle to provide the density required for petabyte-scale continuous HRR datasets. The architectural solution lies in 5D optical storage systems utilizing fused quartz[41].

5D optical storage employs femtosecond laser pulses to induce self-assembled nanogratings within a quartz substrate[41]. These nanogratings induce form birefringence, encoding data across five independent degrees of freedom: three spatial dimensions (multilayer voxel arrays arranged in depth) and two birefringence parameters[41]. The two critical birefringence parameters are the slow-axis direction, denoted as $\psi(x, y)$, and the retardance magnitude, denoted as $\Delta(x, y)$ [41].

Extracting this data at high speeds historically suffered from polarization instability and additive Gaussian noise ($I_{add} \sim N(0, \sigma_{add}^2)$) generated by thermal drift and spectral jitter during laser writing[41]. Furthermore, as the read laser focuses onto a target layer, out-of-focus layers positioned farther from the focal plane intersect a wider beam footprint, exerting a cumulative

distance-dependent modulation on the probe beam[41].
To bypass these noise limitations, modern 5D read-out mechanisms utilize convolutional neural network decoders to solve the polarization detection problem in real-time, effectively isolating the desired layer's birefringence from the cumulative geometric scattering of adjacent layers[41]. This allows the fused quartz to act as a massive, read-only memory bank holding the raw, uncompressed phase maps that will subsequently be formatted into byte streams.

Holographic Data Transport: TIFF and GIF Architectures

The core of the user requirement is the instantiation of the weight as a "byte by byte processable hologram utilizing a tiff-gif spooled at MHz." In this architecture, TIFF and GIF are not static image files; they are highly optimized, standardized payload delivery mechanisms designed to transport spatial phase states from memory directly to the optical plane.

The TIFF Architecture for Continuous Phase Mapping

The Tagged Image File Format (TIFF) is ideally suited for encoding uncompressed, high-dynamic-range holographic data. Unlike traditional rigid array structures, the TIFF specification is a highly extensible container format that deals directly with raw byte alignment and directory offsets[43].
A TIFF file is organized into an 8-byte Image File Header (IFH) and an Image File Directory (IFD)[43]. The IFH immediately declares the byte order—either II for Intel (little-endian) or MM for Motorola (big-endian)[46]. In an FPGA-driven Direct Memory Access (DMA) streaming context, enforcing II little-endian alignment exactly matches the native word alignment of Peripheral Component Interconnect Express (PCIe) engines, dropping byte-swapping translation overhead to absolute zero[46].
The IFD consists of a sequence of 12-byte field entries, each containing a Tag ID, a Field Type, a Length (Count), and a Value Offset[47]. Because the Value Offset can point to any location within the file (even after the image data itself), multiple multi-dimensional holographic layers can be stored contiguously in a single massive stream[46].

TIFF Structure Component	Byte Size	Holographic Streaming Function
Image File Header (IFH)	8 Bytes	Determines Little-Endian DMA alignment (II).
Image File Directory (IFD)	Variable	Contains 12-byte tags indexing the phase maps.
FLOAT / DOUBLE Types	4 / 8 Bytes	Stores ultra-precise continuous phase values

		natively.
Private Tags (>= 32768)	12 Bytes	Out-of-band metadata routing to FPGA schedulers.

For complex weights requiring continuous phase (ϕ) and amplitude (a) precision, the TIFF 6.0 specification natively supports IEEE single-precision (4-byte FLOAT, Type 11) and double-precision (8-byte DOUBLE, Type 12) formats[47]. This allows the raw bytes of a TIFF to represent high-fidelity phase variations necessary for precision D2NNs without interpolation[47]. Furthermore, the architecture leverages modern TIFF extensions. Standard TIFF limits file sizes to 4 GB due to 32-bit offsets[52]. The BigTIFF specification (version 43) utilizes 64-bit offsets, permitting practically unlimited continuous streaming from multi-terabyte quartz storage[52]. Multi-dimensional mapping—such as 8-dimensional HRR conceptual spaces—is natively handled by formats like OME-TIFF and ImageJ hyperstacks, which store multi-dimensional phase matrices contiguously after the first IFD[52]. Finally, the TIFF standard allows for Private Tags (tag numbers 32768 and higher)[54]. The FPGA parser utilizes these private tags (e.g., AOD_RF_CARRIER_FREQ) to extract out-of-band hardware routing instructions alongside the visual hologram, dynamically adjusting the optical hardware without interrupting the data payload[54].

The GIF Data Stream for Continuous LZW Spooling

While TIFF provides uncompressed floating-point precision for continuous phase maps, the Graphics Interchange Format (GIF) excels at rapid, continuous, compressed byte streaming for quantized optical networks. The GIF89a standard utilizes an 8-bit indexed color palette, capping phase states at 256 discrete levels, and relies on lossless Lempel-Ziv-Welch (LZW) compression[57]. For high-speed hardware spooling, the GIF architecture acts as the ultimate pipeline.

The authority on the format is the W3C GIF89a specification[59]. A valid GIF begins with a 6-byte header containing the magic bytes 47 49 46 38 39 61 (GIF89a)[57]. Following the Logical Screen Descriptor and the Global Color Table, the Image Data block dictates the streaming mechanics[57].

The first byte of the Raster Data field is the LZW Minimum Code Size, which indicates the starting bit-width of the LZW dictionary[60]. The dictionary is reset by a Clear code (2^b), and compression terminates with an End-of-Information code ($2^b + 1$)[59]. This explicit dictionary sizing allows an FPGA-based hardware decoder to dynamically adjust its internal shift registers for maximum throughput ahead of the incoming data[59].

The genius of GIF for hardware streaming lies in its sub-block packaging. Instead of presenting a massive monolithic block of compressed data, GIF raster data is packaged into sequential

sub-blocks of 1 to 255 bytes[59]. Each sub-block is prefixed by a single length byte indicating how many data bytes follow[60].

For an FPGA direct memory access (DMA) engine, this structure is revolutionary. The FPGA does not need to buffer an entire multi-megabyte tensor in its limited block RAM; it simply ingests 255-byte chunks continuously, feeding them directly into highly parallelized on-chip LZW decoders[62]. Hardware-accelerated LZW decoders instantiated on modern FPGAs can achieve decompression throughputs exceeding 200 MiB/s per core[64]. A length byte of 0x00 acts as the Block Terminator, seamlessly triggering the end of the frame, followed by a Graphic Control Extension (indicated by 21 F9) to manage phase transition delays, or the trailer byte 3B to terminate the sequence[57]. Furthermore, the Application Extension block (21 FF 0B) allows the stream to specify infinite looping (via the NETSCAPE2.0 identifier), creating a continuously oscillating holographic weight that requires no further instruction from the host CPU[58].

Hardware Modulation: Acousto-Optic Deflectors (AODs) for MHz Spooling

To instantiate the diffractive mask as a physical hologram spooling at MHz frequencies, standard Spatial Light Modulators (SLMs) are entirely insufficient. Conventional liquid crystal SLMs are restricted to the Hz or low kHz ranges due to the physical relaxation times of the liquid crystal molecules[51]. Achieving continuous Megahertz refresh rates necessitates solid-state, inertia-free optical steering elements.

The Physics of Acousto-Optic Deflection

The core modulation hardware for this architecture is the Acousto-Optic Deflector (AOD). An AOD exploits the acousto-optic effect—specifically photoelasticity—to convert radio-frequency (RF) driven acoustic waves into a spatially varying refractive index grating within a transparent crystal medium, typically tellurium dioxide ($\mathrm{TeO_2}$)[68].

When a piezo-transducer bonded to the crystal is driven by electrical signals at high radio frequencies (e.g., 30–350 MHz), it generates a pressure wave that propagates through the medium[70]. This travelling sound wave alters the permittivity of the material, creating a dynamic, volumetric phase mask[69]. The angular deflection θ of the first-order Bragg diffracted beam is governed by the equation:

$$\theta(x,t) = \frac{\lambda}{\nu} f$$

where λ is the optical wavelength, ν is the acoustic velocity in the medium (typically $650\ \mathrm{m/s}$ in $\mathrm{TeO_2}$), and f is the frequency of the driving RF signal[69].

AODs are distinct from Acousto-Optic Modulators (AOMs). While AOMs only modulate the amplitude of the sound wave to alter intensity, AODs adjust both amplitude and frequency

simultaneously, shifting the optical frequency of the deflected beam by an amount equal to the acoustic frequency (creating a Doppler shift) while steering the beam[68]. This allows AODs to achieve microsecond to sub-microsecond switching times, easily sustaining continuous MHz-rate optical refresh rates[68].

Deflector Parameter	Physical Limitation	Architectural Implication for Holography
Response Time	Microsecond to Sub-microsecond	Enables true >1 MHz frame rate spooling.
Access Time	Beam Diameter / Acoustic Velocity	Dictates latency of spatial phase updates.
Resolution (N)	$\Delta f \cdot \tau$ (Bandwidth $\times$ Access Time)	Defines the maximum addressable phase points.
Parasitic Effects	Thermal Lensing, RF Impedance Mismatch	Requires active cooling and precise FPGA calibration.

FPGA RF Synthesis and PCIe Streaming

The AOD does not require a single frequency; it can process multitone RF input signals. By applying multiple frequencies simultaneously, a single AOD can generate highly complex two-dimensional spatial patterns, creating arrays with up to 50x50 distinct beam geometries[73]. Two perpendicularly oriented AODs enable vast 2D diffractive phase mapping[73].

The byte-by-byte spooling pipeline operates via a PCIe 5.0 FPGA accelerator. A PCIe 5.0 x16 link provides a bidirectional raw data rate of 32 GigaTransfers per second (GT/s) per lane, yielding 128 GB/s of total bandwidth[75]. The FPGA utilizes High-Level Synthesis (HLS) to define hardware behavior, enabling direct host-to-FPGA streaming via Direct Memory Access (DMA) using EN_STRM directives[49].

The host memory pumps the GIF sub-blocks or TIFF raw bytes over the PCIe bus into the FPGA's local memory[77]. The on-chip LZW decoders inflate the GIF bytes into an uncompressed array of 8-bit quantized phase levels[59]. This array is instantly mapped to a massive bank of Digital-to-Analog Converters (DACs) capable of GHz sampling rates, which synthesize the multitone RF waveforms required to drive the AOD[68].

By locking the read rate of the PCIe 5.0 bus to the DAC output clock, the entire system operates as a continuous MHz spool. The AOD converts the continuous RF stream into physical sound waves propagating through the $\mathrm{TeO_2}$ crystal. As the laser passes through this

acoustic field, the interference pattern—the neural weight—is instantiated physically, executing the network inference via diffraction at the speed of light[33].

Signal Integrity and Noise Resilience in Low-SNR Regimes

Abandoning exact floating-point mathematics in favor of analog diffractive optics introduces inherent noise, including optical scattering, spatial dispersion, phase quantization errors, and thermal lensing within the AOD[68].

Phase Quantization and Single-Photon Detection

When streaming via GIF, the continuous phase values of the ideal holographic weight are quantized into 256 discrete levels (8-bit palette). In conventional digital systems, drastic quantization severely degrades performance[27]. However, diffractive neural networks exhibit extraordinary resilience to noise[80].

Recent studies in Single-Photon-Detection Neural Networks (SPDNNs) demonstrate that optical architectures maintain high inference accuracy even in extreme low Signal-to-Noise Ratio (SNR) environments[80]. In an SPDNN, optical matrix-vector multiplications occur at ultra-low light intensities governed strictly by Poissonian photon statistics, where the probability of a detector click is $P_{SPD}(\lambda) = 1 - e^{-\lambda}$ [80].

Despite operating with on average less than one detected photon per scalar multiplication (an SNR of roughly 1), these physical networks achieve test accuracies exceeding 90%[80]. This proves that the physical holographic weight does not require 32-bit floating-point precision to function optimally. The stochastic, highly resilient nature of distributed Holographic Reduced Representations ensures that even an imperfectly quantized, noisy GIF stream mapped to an AOD will yield highly accurate cognitive inferences[10].

Adversarial Robustness and Anti-Interference

A critical vulnerability in optical networks is adversarial perturbation in the complex tensor domain. Attack mechanisms, such as the Complex Fast Gradient Sign Method (C-FGSM), have demonstrated that imperceptible physical phase alterations can confuse diffractive models[82].

To counter adversarial interference and hardware-induced scattering, the architecture employs Anti-Interference Diffractive Deep Neural Networks (AI D2NNs)[31]. AI D2NNs utilize task decomposition, splitting the overall network classification task across multiple parallel spatial channels[30]. By spooling distinct TIFF or GIF streams to separate, parallel AODs, the optimization difficulty of individual phase masks is reduced[30]. If noise or adversarial interference corrupts one spatial channel, the ensemble inference remains robust, maintaining high accuracy even in dynamic, multi-object scenarios without requiring error-correction overhead in the byte stream[27].

Synthesis and Conclusion

The architecture detailed in this report dismantles the prevailing doctrine of modern artificial intelligence hardware. It forcibly transitions the field away from discrete, localized shards of

mathematical relationships constrained by the von Neumann bottleneck, advancing instead toward continuous, distributed Vector Symbolic Architectures (VSAs) reliant upon physical optical convolution[5].

By framing the neural weight as a byte-by-byte processable physical hologram, the computational bottleneck shifts entirely from arithmetic calculation to I/O bandwidth. Spooling 8-bit quantized lossless LZW-encoded GIF streams or raw, high-dynamic-range uncompressed TIFF formats leverages the inherent structural efficiencies of these protocols[43]. These byte streams—transported via 128 GB/s PCIe 5.0 Direct Memory Access pipelines to FPGA-driven RF synthesizers—excite Acousto-Optic Deflectors, creating fully reconfigurable, volumetric diffractive arrays updating at multiple Megahertz[70].

When combined with the inherent $O(1)$ circular convolution capabilities of $4f$ Fourier optics, the network performs ultra-high-dimensional cognitive associations instantaneously[12]. The system operates efficiently within low-SNR, Poissonian light limits, natively resisting quantization degradation[80]. The ultimate outcome is a continuously streaming optical intelligence engine—one that abandons mathematical shards for physical light fields, computing holistically at the speed of light, and scaling orders of magnitude beyond current electronic paradigms.

Works cited

1. Noninvasive Monitoring of Blood Glucose Using Color-Coded ... - PMC, https://pmc.ncbi.nlm.nih.gov/articles/PMC6232738/
2. Holographic associative memory, https://en-academic.com/dic.nsf/enwiki/2559372
3. UMI - viXra.org, https://vixra.org/pdf/1712.0136v1.pdf
4. Full text of "Technological Singularity Wikibook" - Internet Archive, https://archive.org/stream/TechnologicalSingularityWikibook/TechnologicalSingularityWikibook_djvu.txt
5. Holographic functions and neural networks - arXiv, https://arxiv.org/pdf/2605.22666
6. NEURAL NETWORKS AS A BASIS FOR QUANTUM ASSOCIATIVE, https://www.semanticscholar.org/paper/NEURAL-NETWORKS-AS-A-BASIS-FOR-QUANTUM-ASSOCIATIVE-Perus/f8deca83367163b78d8f2cfb433d408060e8b995
7. [2605.22666] Holographic functions and neural networks - arXiv, https://arxiv.org/abs/2605.22666
8. Holographic Reduced Representations for Working Memory, https://jewlscholar.mtsu.edu/bitstreams/32f0ed2f-63dd-4478-96a6-8a1941ce15c1/download
9. VSA- and HRR-related publications - Tony Plate's Publications, http://d-reps.org/vsapubs.html
10. Associative Long Short-Term Memory, http://proceedings.mlr.press/v48/danihelka16.pdf
11. (PDF) Learning Holographic Reduced Representations with Clifford,

https://www.researchgate.net/publication/414686640_Learning_Holographic_Reduced_Representations_with_Clifford_Variational_Autoencoders/download
12. Modelling neural probabilistic computation using vector symbolic, https://pmc.ncbi.nlm.nih.gov/articles/PMC11655797/
13. Bridging Cognitive Architectures and Generative Models with Vector, https://ojs.aaai.org/index.php/AAAI-SS/article/download/27686/27459/31737
14. Iterative projection meets sparsity regularization, https://www.light-am.com/article/pdf/preview/LAM2022050014.pdf
15. Real-time Optical Realization of Circle Hough Transform with, https://www.researchgate.net/publication/301422317_Real-time_Optical_Realization_of_Circle_Hough_Transform_with_Incoherent_Light
16. Solution to the issue of high-order diffraction images for cylindrical, https://opg.optica.org/oe/fulltext.cfm?uri=oe-32-9-14978
17. Tony Plate's Publications, http://www.d-reps.org/
18. Using Diffraction Deep Neural Networks for Indirect Phase Recovery, https://pmc.ncbi.nlm.nih.gov/articles/PMC10819540/
19. Generative Adversarial Optical Networks Using Diffractive Layers for, https://www.mdpi.com/2304-6732/13/1/94
20. Parallel photonic acceleration processor for matrix–matrix, https://opg.optica.org/ol/fulltext.cfm?uri=ol-48-12-3231
21. Photonic matrix multiplication lights up photonic accelerator and, https://www.researchgate.net/publication/358351958_Photonic_matrix_multiplication_lights_up_photonic_accelerator_and_beyond
22. The MFT System Part 1: Silicon Photonics - Optalysys, https://optalysys.com/resource/the-mft-system-part-1-silicon-photonics/
23. Matrix processing with nanophotonics - Lightmatter - Medium, https://medium.com/lightmatter/matrix-processing-with-nanophotonics-998e294dabc1
24. Photonic matrix multiplication lights up photonic accelerator and, https://research.cuhk.edu.hk/files/256400995/388371591.pdf
25. Photonic Matrix Computing: From Fundamentals to Applications - PMC, https://pmc.ncbi.nlm.nih.gov/articles/PMC8308143/
26. Diffractive optical elements 75 years on: from micro ... - Researching, https://www.researching.cn/articles/OJ2130a292345c96aa
27. A Residual Optronic Convolutional Neural Network for SAR Target, https://www.mdpi.com/2304-6732/12/7/678
28. Metasurface-Based Image Classification Using Diffractive Deep, https://www.mdpi.com/2079-4991/14/22/1812
29. Study of simulation model of diffractive neural network in liquid, https://epjam.edp-open.org/articles/epjam/full_html/2026/01/epjam250025/epjam250025.html
30. Simulated task-decomposed multi-channel diffractive neural, https://opg.optica.org/oe/abstract.cfm?uri=oe-34-18-33357
31. Anti-interference diffractive deep neural networks for multi-object, https://pmc.ncbi.nlm.nih.gov/articles/PMC12868834/

32. Ensemble learning of diffractive optical networks, https://www.light-am.com/article/pdf/preview/xjzz-2021-10-123.pdf
33. Optical Phase Space Measurements and Applications ... - UC Berkeley, https://escholarship.org/content/qt5kk7w66s/qt5kk7w66s.pdf
34. Compact and effective zero-order noise suppression in Fresnel, https://www.spiedigitallibrary.org/conference-proceedings-of-spie/13573/1357305/Compact-and-effective-zero-order-noise-suppression-in-Fresnel-holography/10.1117/12.3063744.full
35. Optical Diffraction Tomography: A Resolution Study - AMS Laurea, https://amslaurea.unibo.it/id/eprint/21258/1/odt.pdf
36. Physics-guided deep unfolding network for snapshot 3D imaging, https://www.researching.cn/ArticlePdf/m00105/2026/5/5/056024.pdf
37. Optical Diffractive Convolutional Neural Networks Implemented in, https://pmc.ncbi.nlm.nih.gov/articles/PMC10302459/
38. Rotation Invariant Parallel Signal Processing Using a Diffractive, https://www.mdpi.com/2076-3417/12/1/439
39. Pixel Super-Resolution Phase Retrieval for Lensless On-Chip ... - PMC, https://pmc.ncbi.nlm.nih.gov/articles/PMC9265759/
40. A method to improve the computational performance of nonlinear all, https://www.tandfonline.com/doi/full/10.1080/15599612.2023.2209624
41. Multi-layer 5D Optical Data Storage - arXiv, https://arxiv.org/pdf/2508.20106
42. Chinese Scientists Built A 3D Holographic Storage System ... - Reddit, https://www.reddit.com/r/InterstellarKinetics/comments/1s4bq22/exclusive_chinese_scientists_built_a_3d/
43. TIFF File Format Summary - Paul Bourke, https://paulbourke.net/dataformats/tiff/tiff_summary.pdf
44. OGC GeoTIFF Standard - OGC Standards, https://docs.ogc.org/is/19-008r4/19-008r4.html
45. The Hidden Art of Crafting TIFF Files: How to Create TIFF Image, https://files.crisesnotes.com/article/the-hidden-art-of-crafting-tiff-files-how-to-create-tiff-image-files-like-a-pro
46. Revision 6.0 - AlternaTIFF, https://www.alternatiff.com/resources/TIFF6.pdf
47. TIFF - Revision 6.0 - Publications Office of the EU, https://op.europa.eu/documents/3938058/5896514/TIFF-v6/9340926d-6a7a-63c9-5fd7-21c435dcdfeb
48. Description of Exif file format - MIT Media Lab, https://www.media.mit.edu/pia/Research/deepview/exif.html
49. UCLA Electronic Theses and Dissertations - eScholarship.org, https://escholarship.org/content/qt072909j7/qt072909j7.pdf
50. TIFF Revision 5.0, https://cool.culturalheritage.org/bytopic/imaging/std/tiff5.html
51. Optical Reservoir Computing using multiple light scattering ... - arXiv, https://arxiv.org/html/1907.00657v2
52. tifffile | Documentation - Christoph Gohlke, https://www.cgohlke.com/docs/tifffile/
53. Loading and Saving Large TIFF/BigTIFF Files - LEADTOOLS,

https://www.leadtools.com/help/sdk/main/api/loading-and-saving-large-tiff-bigtiff-files.html
54. Custom tags on Tiff Image: Private Tags - c++ - Stack Overflow, https://stackoverflow.com/questions/42954220/custom-tags-on-tiff-image-private-tags
55. ANSI X9.100-181-2014 (R2021) TIFF Image Format for Image, https://x9.org/wp-content/uploads/2021/01/ANSI-X9.100-181-2014-R2021.pdf
56. Defining New TIFF Tags - LibTIFF, http://libtiff.maptools.org/addingtags.html
57. GIF - Wikipedia, https://en.wikipedia.org/wiki/GIF
58. What Is a GIF File? Complete Format Guide - CoolUtils, https://www.coolutils.com/blog/what-is-gif-file/
59. GIF file extension - Graphics Interchange Format File, https://file-extensions.com/docs/gif
60. What's In A GIF - Bits and Bytes - GIFLib, https://giflib.sourceforge.net/whatsinagif/bits_and_bytes.html
61. Image file type and format guide - Media - MDN Web Docs - Mozilla, https://developer.mozilla.org/en-US/docs/Web/Media/Guides/Formats/Image_types
62. Network-conscious GIF image transmission over the Internet, https://sites.cs.ucsb.edu/~pconrad/files/paper014.pdf
63. 3MF Project: What's In A GIF - Bit by Byte - Matthew Flickinger, https://www.matthewflickinger.com/lab/whatsinagif/bits_and_bytes.asp
64. redwarp/lzw: Fast implementation of the LZW encoder and decoder, https://github.com/redwarp/lzw
65. Decompressor for hardware applications, https://aait.od.ua/index.php/journal/article/download/115/115
66. The Tiniest GIF Ever - Probably Programming, http://probablyprogramming.com/2009/03/15/the-tiniest-gif-ever
67. 4F Optical Neural Network Acceleration: An Architectural Perspective, https://nanocad.ee.ucla.edu/wp-content/papercite-data/pdf/ip18.pdf
68. What is Acousto-optic deflector? Meaning, Examples, Use Cases, https://quantumopsschool.com/blog/acousto-optic-deflector/
69. Acousto-optics - Wikipedia, https://en.wikipedia.org/wiki/Acousto-optics
70. Acousto-optic deflectors in experimental neuroscience, https://diposit.ub.edu/server/api/core/bitstreams/417ed803-6b23-455a-9adb-7b39971bd1dc/content
71. AA Opto-Electronic - Acousto-optic devices and RF drivers, https://aaoptoelectronic.com/
72. Megahertz Spatial Light Modulator Achieves Rapid, https://quantumzeitgeist.com/megahertz-spatial-light-modulator-achieves/
73. Model-Based Real-Time Synthesis of Acousto-Optically Generated, https://arxiv.org/pdf/2512.16774
74. Ultrahigh-throughput single-pixel complex-field microscopy with, https://pmc.ncbi.nlm.nih.gov/articles/PMC12340123/
75. PCIe 5.0 vs 4.0: A Comprehensive Technical Deep Dive for Engineers,

https://www.wevolver.com/article/pcie-50-vs-40-a-comprehensive-technical-deep-dive-for-engineers
76. PCIe 4.0 vs PCIe 5.0 Graphics Cards | Gaming Guide 2026 - Box, https://box.co.uk/blog/pcie-4-vs-pcie-5-graphics-cards-gaming-performance
77. Do OS abstractions make sense on FPGAs? - USENIX, https://www.usenix.org/system/files/osdi20-korolija.pdf
78. Optical next generation reservoir computing - arXiv, https://arxiv.org/html/2404.07857v3
79. Optoelectronic nonlinear Softmax operator based on diffractive, https://opg.optica.org/oe/fulltext.cfm?uri=oe-32-15-26458
80. Quantum-limited stochastic optical neural networks operating at a, https://pmc.ncbi.nlm.nih.gov/articles/PMC11698857/
81. An optical neural network using less than 1 photon per multiplication, https://pmc.ncbi.nlm.nih.gov/articles/PMC8748769/
82. Physical Adversarial Attacks of Diffractive Deep Neural Networks, https://ycunxi.github.io/cunxiyu/papers/lbr_dac21_advonn.pdf